

```
In [0]: from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
        import warnings
        warnings.filterwarnings("ignore")
```

Plotting code snippet below :

```
In [0]: %matplotlib inline
        import matplotlib.pyplot as plt
        import numpy as np
        import time

        def plt_dynamic (x , vy , ty , cx , colors = ['b']):
            ax.plot (x , vy , 'b' , label = "Validation Loss")
            ax.plot (x , vy , 'r' , label = "Train Loss")
            plt.legend()
            plt.grid()
            fig.canvas.draw()
```

Loading the train and test data

```
In [3]: #shuffled and split between train and test sets

        (X_train , y_train) , (X_test , y_test) = mnist.load_data();
```

```
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [=====] - 3s 0us/step
```

Printing the training examples

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [7]: print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0		
0	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	25	
5	247	127	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	15	
4	170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0		
0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	8	
2	82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	25
3	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	15	
4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	24
1	225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0		
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
0	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	18	
7																			

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 25
3 253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 25
3 253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 19
5 80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 1
6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]

```

0 means white 255 means black Between 0 and 255 means grey

```
X_train = X_train/255
X_test = X_test/255
```

Observation:

Normalization being done here

```
print(X_train[0])
```

[illegible]

0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686

0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

```
In [10]: print("Class label of first image :", y_train[0])
```

```
#converting class labels into one hot encoding
```

```
Y_train = np_utils.to_categorical(y_train, 10)
```

```
Y_test = np_utils.to_categorical(y_test, 10)
```

```
print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```
In [0]: from keras.models import Sequential
```



```
from keras.layers import Dense, Activation
```

```
In [0]: output_dim = 10  
input_dim = X_train.shape[1]  
  
batch_size = 128  
nb_epoch = 20
```

```
In [0]: model = Sequential()  
  
model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

MLP + SGD OPTIMIZER + SIGMOID ACTIVATION

```
In [14]: model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics  
=['accuracy'])
```

```
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_  
epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 26us/step - loss: 1.3
017 - acc: 0.6960 - val_loss: 0.8110 - val_acc: 0.8389

Epoch 2/20

60000/60000 [=====] - 1s 19us/step - loss: 0.7
186 - acc: 0.8418 - val_loss: 0.6070 - val_acc: 0.8639

Epoch 3/20

60000/60000 [=====] - 1s 19us/step - loss: 0.5
884 - acc: 0.8593 - val_loss: 0.5251 - val_acc: 0.8747

Epoch 4/20

```
60000/60000 [=====] - 1s 19us/step - loss: 0.5
262 - acc: 0.8687 - val_loss: 0.4795 - val_acc: 0.8822
Epoch 5/20
60000/60000 [=====] - 1s 20us/step - loss: 0.4
883 - acc: 0.8756 - val_loss: 0.4499 - val_acc: 0.8868
Epoch 6/20
60000/60000 [=====] - 1s 21us/step - loss: 0.4
622 - acc: 0.8808 - val_loss: 0.4282 - val_acc: 0.8903
Epoch 7/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4
429 - acc: 0.8837 - val_loss: 0.4124 - val_acc: 0.8924
Epoch 8/20
60000/60000 [=====] - 1s 21us/step - loss: 0.4
279 - acc: 0.8871 - val_loss: 0.3995 - val_acc: 0.8942
Epoch 9/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4
158 - acc: 0.8895 - val_loss: 0.3893 - val_acc: 0.8970
Epoch 10/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4
058 - acc: 0.8916 - val_loss: 0.3805 - val_acc: 0.8989
Epoch 11/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3
973 - acc: 0.8933 - val_loss: 0.3732 - val_acc: 0.9005
Epoch 12/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3
900 - acc: 0.8948 - val_loss: 0.3670 - val_acc: 0.9016
Epoch 13/20
60000/60000 [=====] - 1s 18us/step - loss: 0.3
836 - acc: 0.8968 - val_loss: 0.3616 - val_acc: 0.9024
Epoch 14/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3
780 - acc: 0.8977 - val_loss: 0.3564 - val_acc: 0.9034
Epoch 15/20
60000/60000 [=====] - 1s 21us/step - loss: 0.3
729 - acc: 0.8986 - val_loss: 0.3518 - val_acc: 0.9044
Epoch 16/20
60000/60000 [=====] - 1s 21us/step - loss: 0.3
684 - acc: 0.8997 - val_loss: 0.3481 - val_acc: 0.9052
Epoch 17/20
```

```

60000/60000 [=====] - 1s 21us/step - loss: 0.3
643 - acc: 0.9004 - val_loss: 0.3445 - val_acc: 0.9066
Epoch 18/20
60000/60000 [=====] - 1s 20us/step - loss: 0.3
606 - acc: 0.9011 - val_loss: 0.3414 - val_acc: 0.9074
Epoch 19/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3
571 - acc: 0.9018 - val_loss: 0.3382 - val_acc: 0.9088
Epoch 20/20
60000/60000 [=====] - 1s 19us/step - loss: 0.3
540 - acc: 0.9023 - val_loss: 0.3355 - val_acc: 0.9089

```

```

In [15]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']

```

```
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.3354729148030281

Test accuracy: 0.9089



```
In [16]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_d
im,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

```
In [17]: model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy',  
    metrics=['accuracy'])
```

```
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,  
    verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 77us/step - loss: 2.2

696 - acc: 0.2015 - val_loss: 2.2274 - val_acc: 0.3192

Epoch 2/20

60000/60000 [=====] - 4s 74us/step - loss: 2.1

849 - acc: 0.4129 - val_loss: 2.1323 - val_acc: 0.4897

Epoch 3/20

60000/60000 [=====] - 4s 74us/step - loss: 2.0

732 - acc: 0.5437 - val_loss: 1.9958 - val_acc: 0.5616

Epoch 4/20

60000/60000 [=====] - 4s 73us/step - loss: 1.9

136 - acc: 0.6075 - val_loss: 1.8044 - val_acc: 0.6549

Epoch 5/20

60000/60000 [=====] - 4s 75us/step - loss: 1.7

052 - acc: 0.6553 - val_loss: 1.5760 - val_acc: 0.7009

Epoch 6/20

60000/60000 [=====] - 4s 74us/step - loss: 1.4

768 - acc: 0.6995 - val_loss: 1.3489 - val_acc: 0.7233

Epoch 7/20

60000/60000 [=====] - 4s 74us/step - loss: 1.2

680 - acc: 0.7345 - val_loss: 1.1597 - val_acc: 0.7576

Epoch 8/20

60000/60000 [=====] - 4s 73us/step - loss: 1.1

004 - acc: 0.7637 - val_loss: 1.0120 - val_acc: 0.7852

Epoch 9/20

60000/60000 [=====] - 4s 74us/step - loss: 0.9

716 - acc: 0.7850 - val_loss: 0.8998 - val_acc: 0.8035

```

Epoch 10/20
60000/60000 [=====] - 4s 74us/step - loss: 0.8
720 - acc: 0.8021 - val_loss: 0.8133 - val_acc: 0.8171
Epoch 11/20
60000/60000 [=====] - 4s 75us/step - loss: 0.7
937 - acc: 0.8145 - val_loss: 0.7431 - val_acc: 0.8244
Epoch 12/20
60000/60000 [=====] - 4s 73us/step - loss: 0.7
306 - acc: 0.8256 - val_loss: 0.6868 - val_acc: 0.8373
Epoch 13/20
60000/60000 [=====] - 4s 72us/step - loss: 0.6
793 - acc: 0.8341 - val_loss: 0.6406 - val_acc: 0.8432
Epoch 14/20
60000/60000 [=====] - 4s 73us/step - loss: 0.6
373 - acc: 0.8411 - val_loss: 0.6022 - val_acc: 0.8485
Epoch 15/20
60000/60000 [=====] - 4s 74us/step - loss: 0.6
022 - acc: 0.8476 - val_loss: 0.5700 - val_acc: 0.8553
Epoch 16/20
60000/60000 [=====] - 4s 73us/step - loss: 0.5
726 - acc: 0.8531 - val_loss: 0.5430 - val_acc: 0.8614
Epoch 17/20
60000/60000 [=====] - 4s 73us/step - loss: 0.5
475 - acc: 0.8580 - val_loss: 0.5198 - val_acc: 0.8646
Epoch 18/20
60000/60000 [=====] - 4s 74us/step - loss: 0.5
258 - acc: 0.8621 - val_loss: 0.4999 - val_acc: 0.8681
Epoch 19/20
60000/60000 [=====] - 4s 72us/step - loss: 0.5
069 - acc: 0.8665 - val_loss: 0.4821 - val_acc: 0.8723
Epoch 20/20
60000/60000 [=====] - 4s 75us/step - loss: 0.4
902 - acc: 0.8698 - val_loss: 0.4675 - val_acc: 0.8753

```

```

In [27]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ADAM OPTIMIZER +SIGMOID

ACTIVATION

```
In [19]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy',
, metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
=====	=====	=====

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 97us/step - loss: 0.5567 - acc: 0.8541 - val_loss: 0.2564 - val_acc: 0.9262
Epoch 2/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2213 - acc: 0.9348 - val_loss: 0.1892 - val_acc: 0.9432
Epoch 3/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1625 - acc: 0.9519 - val_loss: 0.1414 - val_acc: 0.9565


```
Epoch 4/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1
246 - acc: 0.9633 - val_loss: 0.1230 - val_acc: 0.9608
Epoch 5/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
994 - acc: 0.9709 - val_loss: 0.0976 - val_acc: 0.9705
Epoch 6/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
785 - acc: 0.9767 - val_loss: 0.0903 - val_acc: 0.9718
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
640 - acc: 0.9806 - val_loss: 0.0847 - val_acc: 0.9726
Epoch 8/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
512 - acc: 0.9848 - val_loss: 0.0774 - val_acc: 0.9751
Epoch 9/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
412 - acc: 0.9881 - val_loss: 0.0702 - val_acc: 0.9780
Epoch 10/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
332 - acc: 0.9906 - val_loss: 0.0667 - val_acc: 0.9793
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
275 - acc: 0.9923 - val_loss: 0.0636 - val_acc: 0.9808
Epoch 12/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
211 - acc: 0.9945 - val_loss: 0.0628 - val_acc: 0.9815
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0
174 - acc: 0.9957 - val_loss: 0.0606 - val_acc: 0.9823
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0
138 - acc: 0.9966 - val_loss: 0.0618 - val_acc: 0.9819
Epoch 15/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0
112 - acc: 0.9974 - val_loss: 0.0612 - val_acc: 0.9827
Epoch 16/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
088 - acc: 0.9980 - val_loss: 0.0739 - val_acc: 0.9798
```

```

Epoch 17/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
074 - acc: 0.9984 - val_loss: 0.0650 - val_acc: 0.9832
Epoch 18/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
055 - acc: 0.9990 - val_loss: 0.0732 - val_acc: 0.9813
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
045 - acc: 0.9991 - val_loss: 0.0679 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
041 - acc: 0.9990 - val_loss: 0.0704 - val_acc: 0.9814

```

```

In [20]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epo
chs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter vali
dation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal
to number of epochs


vy = history.history['val_loss']

```

```
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)  
plt.show()
```

Test score: 0.07042204804238318

Test accuracy: 0.9814




```
In [28]: w_after = model_sigmoid.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 3, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w,color='b')  
plt.xlabel('Hidden Layer 1')  
  
plt.subplot(1, 3, 2)  
plt.title("Trained model Weights")
```

```

ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + RELU ACTIVATION + SGD OPTIMIZER

```

In [22]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
In [23]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 76us/step - loss: 0.7554 - acc: 0.7889 - val_loss: 0.3819 - val_acc: 0.8929

Epoch 2/20

60000/60000 [=====] - 4s 75us/step - loss: 0.3461 - acc: 0.9019 - val_loss: 0.2969 - val_acc: 0.9148

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2848 - acc: 0.9193 - val_loss: 0.2581 - val_acc: 0.9271

Epoch 4/20

60000/60000 [=====] - 4s 73us/step - loss: 0.2520 - acc: 0.9285 - val_loss: 0.2377 - val_acc: 0.9338

Epoch 5/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2292 - acc: 0.9355 - val_loss: 0.2193 - val_acc: 0.9374

Epoch 6/20

60000/60000 [=====] - 4s 72us/step - loss: 0.2118 - acc: 0.9405 - val_loss: 0.2051 - val_acc: 0.9419

Epoch 7/20

```
60000/60000 [=====] - 4s 72us/step - loss: 0.1
972 - acc: 0.9449 - val_loss: 0.1937 - val_acc: 0.9462
Epoch 8/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
851 - acc: 0.9481 - val_loss: 0.1829 - val_acc: 0.9477
Epoch 9/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1
746 - acc: 0.9508 - val_loss: 0.1757 - val_acc: 0.9496
Epoch 10/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
651 - acc: 0.9540 - val_loss: 0.1670 - val_acc: 0.9519
Epoch 11/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
570 - acc: 0.9563 - val_loss: 0.1603 - val_acc: 0.9537
Epoch 12/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1
493 - acc: 0.9585 - val_loss: 0.1545 - val_acc: 0.9553
Epoch 13/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
424 - acc: 0.9603 - val_loss: 0.1524 - val_acc: 0.9555
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1
362 - acc: 0.9622 - val_loss: 0.1450 - val_acc: 0.9583
Epoch 15/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
305 - acc: 0.9636 - val_loss: 0.1420 - val_acc: 0.9590
Epoch 16/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
253 - acc: 0.9650 - val_loss: 0.1370 - val_acc: 0.9607
Epoch 17/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1
203 - acc: 0.9666 - val_loss: 0.1335 - val_acc: 0.9614
Epoch 18/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
160 - acc: 0.9675 - val_loss: 0.1291 - val_acc: 0.9639
Epoch 19/20
60000/60000 [=====] - 4s 72us/step - loss: 0.1
117 - acc: 0.9687 - val_loss: 0.1257 - val_acc: 0.9636
Epoch 20/20
```

```
60000/60000 [=====] - 4s 71us/step - loss: 0.1077 - acc: 0.9699 - val_loss: 0.1222 - val_acc: 0.9652
```

```
In [24]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.12223365796022118
```

```
Test accuracy: 0.9652
```



```
In [29]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```



```
plt.xlabel('Output Layer ' )  
plt.show()
```

